

Lab3: Image Stitching

Teaching assistant: Abhinav
Rai
e0403959@u.nus.edu

Objective

The goal of this assignment is to do image stitching – how multiple images can form a panorama. To do this, you would need to implement feature detectors, feature descriptors, keypoints matching and homography computation algorithms .

1 Task 1 Keypoint Detection (10 points)

1.1 Task 1 Description

- For this task we will implement the Harris Corner Detector

1.2 Steps Task 1

- Compute Image Gradient I_x and I_y at each point in the image.
- Compute the Hessian Matrix H for each window.

$$H = \sum_{(x,y) \in W} w_{x,y} \begin{bmatrix} I_x^2 & I_x I_y \\ I_x I_y & I_y^2 \end{bmatrix} \quad (1)$$

- Compute Corner Response.

$$R = \text{Det}(H) - k (\text{Trace}(H))^2 \quad (2)$$



- Find the points whose surrounding windows has a larger corner response ($R \geq \text{threshold}$). Take the points of local maxima i.e., perform Non-Maximum Supression. This again has already been implemented.

1.3 Hints

- Use `skimage.filters.sobel_v / sobel_h`
- Use `scipy.ndimage.filters.convolve`

1.4 Expected Result Task 1


Detected Corners on letter boxDetected Corners on checker



2 Task 2a Keypoint Description (14 points)

2.1 Task 2a Description

- Implement a simple descriptor that describes each keypoint with normalized intensity in a patch (e.g. size of 5*5 pixels) around it.



- Given the Feature F , perform normalization: $F = \frac{F - \mu}{\sigma}$ if $\sigma > 0$ else $F - \mu$

3 Task 2b Keypoint Matching (14 points)

3.1 Task 2b Description

- Calculate Euclidean Distance between all pairs of descriptors from image 1 and image 2.
- If the distance to the closest vector is significantly (by a given factor) smaller than the distance to the second-closest, we consider it a match. The output of the function is an array where each row holds the indices of one pair of matching descriptors.

3.2 Hints

- Use `scipy.spatial.distance.cdist()`

3.3 Expected Results Task 2



4 Task 3 Homography Estimation (18 points)

4.1 Task 3 Description

- For this task we will be using DLT algorithm to compute the homography matrix. Specifically, we will use the normalized DLT algorithm.

4.2 Hints

- Use `np.linalg.svd()`
- Use `transform_homography` function in `image_stitching.py`
- You will not get accurate stitching results in this step, as there are a lot of erroneous matches which need to be filtered. We will implement RANSAC in the next task for that.

5 Task 4 RANSAC for Descriptor Matching (18 points)


5.1 Task 4 Description

- RANSAC is an algorithm that is used to exclude outliers, by iteratively finding out a set of good fits and then using the best fit to do keypoint matching.

5.2 Steps Task 4

- Select a random set of `n_samples` of matches.
- Compute Homography Matrix - Refer to `def compute_homography(p1, p2)`
- Find inliers using the provided threshold - Compute Sum of Squared Difference (*SSD*) on transformed `M2` (*i.e. matched2*) using homography and `M1` (*i.e. matched1*), and pick those lower than the given threshold as inliers.
- Repeat the above steps for `n_iters` and keep the largest set of inliers.
- Recompute the least squares estimate using only the inliers.

5.3 Expected Results Task 4 (22 points)




5.4 Hints

- Use `compute_homography` and `transform_homography`
- You will need to tune the parameters `sampling_ratio`, `n_iters` and `threshold` to generate accurate results.
- The final result may vary across iterations because of random sampling at each step.

6 Task 5 Scale Invariant Feature Transform(SIFT) (18 points)

6.1 Task 5 Description

- Implement a simplified version of SIFT descriptor.
- For each keypoint, take a 16×16 patch and divide it into 4×4 grid cells each of length 4.
- For illustration 8×8 patch and 2×2 grid of cells are shown below.



- Each cell should have a histogram of the local distribution of gradients in 8 orientations.
- Appending these histograms together will give you a $4*4*8 = 128$ -dimensional vector.
- For an image sample, the gradient magnitude m and orientation θ are computed using pixel differences. Also, features should be normalized to unit length.

$$m(x, y) = \sqrt{(L(x+1, y) - L(x-1, y))^2 + (L(x, y+1) - L(x, y-1))^2}$$

$$\theta(x, y) = \text{atan2}(L(x, y+1) - L(x, y-1), L(x+1, y) - L(x-1, y))$$

6.2 Expected Results Task 5

